

Boosting

68_Adnan Shaikh

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: import pandas as pd
df=pd.read_csv('./apples_and_oranges.csv')
```

```
In [3]: df.head()
```

```
Out[3]:
```

	Weight	Size	Class
0	69	4.39	orange
1	69	4.21	orange
2	65	4.09	orange
3	72	5.85	apple
4	67	4.70	orange

```
In [4]: df.shape
```

```
Out[4]: (40, 3)
```

```
In [5]: x=df.drop("Class",axis="columns")
y=df.Class
```

```
In [6]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=1)
```

```
In [7]: x_test.shape
```

```
Out[7]: (8, 2)
```

```
In [8]: x_train.shape
```

```
Out[8]: (32, 2)
```

```
In [9]: from sklearn.ensemble import AdaBoostClassifier
```

```
In [10]: ada=AdaBoostClassifier(n_estimators=100,base_estimator=None,learning_rate=1,random_state=1)
ada.fit(x_train,y_train)
```

```
Out[10]: AdaBoostClassifier(learning_rate=1, n_estimators=100, random_state=1)
```

```
In [11]: y_pred=ada.predict(x_test)
```

```
In [12]: from sklearn.metrics import confusion_matrix
cm=confusion_matrix(y_test,y_pred)
print(cm)
```

```
[[3 0]
 [0 5]]
```

```
In [13]: accuracy=float(cm.diagonal().sum())/len(y_test)
```

```
In [14]: print("Accuracy of Adaboost for Given Data set :",accuracy)
```

Accuracy of Adaboost for Given Data set : 1.0

```
In [15]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
apple	1.00	1.00	1.00	3
orange	1.00	1.00	1.00	5
accuracy			1.00	8
macro avg	1.00	1.00	1.00	8
weighted avg	1.00	1.00	1.00	8

Gradient Boosting

```
In [16]: from sklearn.ensemble import GradientBoostingClassifier
```

```
In [17]: gb = GradientBoostingClassifier(n_estimators=100)
```

```
gb.fit(x_train,y_train)
```

```
Out[17]: GradientBoostingClassifier()
```

```
In [18]: y_pred = gb.predict(x_test)
```

```
In [19]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
apple	1.00	1.00	1.00	3
orange	1.00	1.00	1.00	5
accuracy			1.00	8
macro avg	1.00	1.00	1.00	8
weighted avg	1.00	1.00	1.00	8

Xtreme Gradient Boosting

```
In [20]: from xgboost import XGBClassifier
```

```
In [21]: xgb = XGBClassifier(n_estimators=200,reg_alpha=1)
```

```
In [22]: xgb.fit(x_train,y_train)
```

[15:02:49] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

```
Out[22]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
      colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
      importance_type='gain', interaction_constraints='',
      learning_rate=0.300000012, max_delta_step=0, max_depth=6,
      min_child_weight=1, missing=nan, monotone_constraints='()',
      n_estimators=200, n_jobs=12, num_parallel_tree=1, random_state=0,
      reg_alpha=1, reg_lambda=1, scale_pos_weight=1, subsample=1,
      tree_method='exact', validate_parameters=1, verbosity=None)
```

```
In [23]: y_pred = xgb.predict(x_test)
```

```
In [24]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
apple	0.60	1.00	0.75	3
orange	1.00	0.60	0.75	5
accuracy			0.75	8
macro avg	0.80	0.80	0.75	8
weighted avg	0.85	0.75	0.75	8